

`github.com/ernestyalumni/advanced-tensorflow`

# Advanced TensorFlow

The Datasets API (and on the GPU(s))

Ernest Yeung


<[ernestyalumni@gmail.com](mailto:ernestyalumni@gmail.com)>

Twitter, IG, github: @ernestyalumni

## GPU vs. CPU



*If you were plowing a field, which would you rather use: Two strong oxen or 1024 chickens?*  
--Attributed to Seymour Cray




## Intel Xeon Phi 7210

- 64 cores
- 215 W
- \$1881.00

## NVIDIA GeForce GTX 1080Ti

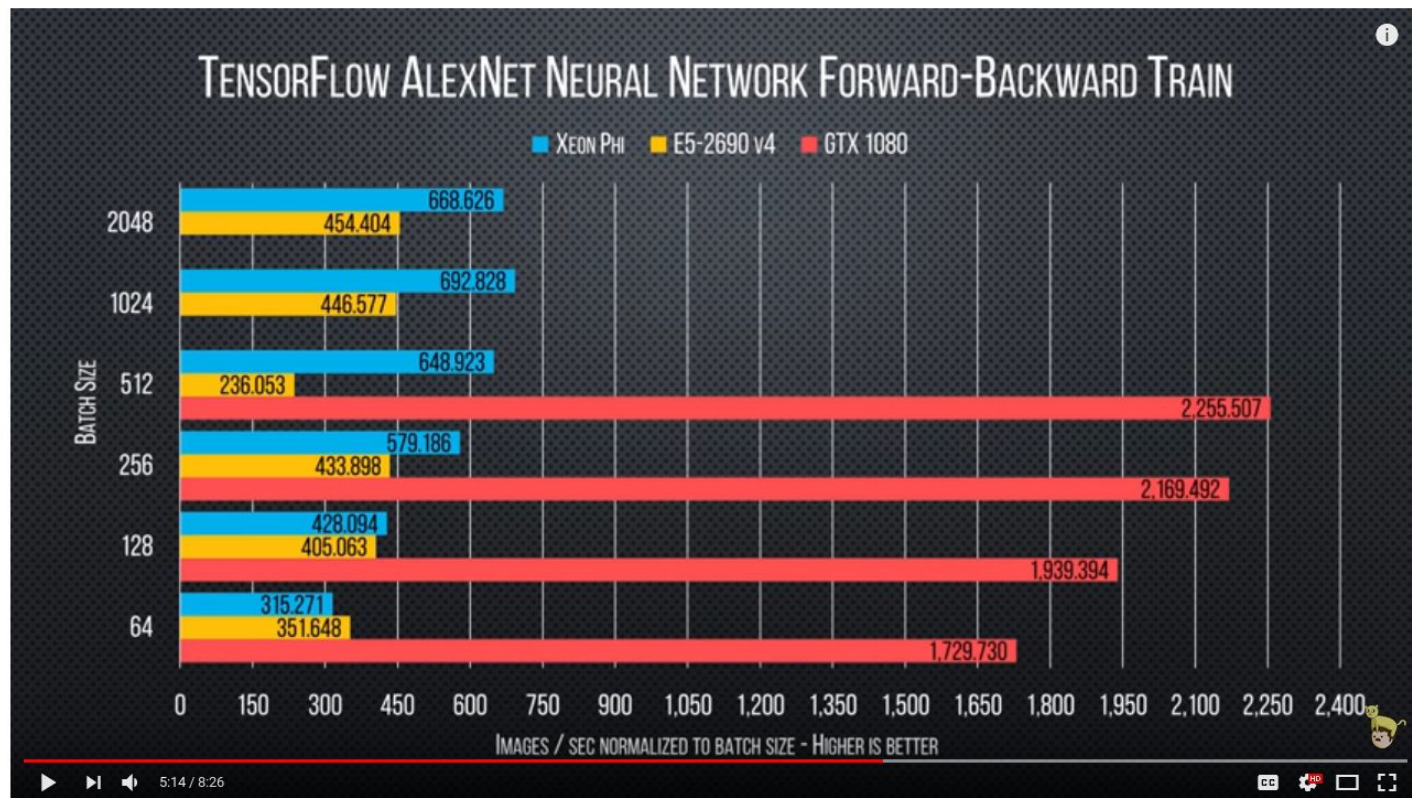
- 3584 CUDA cores
- 250 W
- \$769





# Linus of Linus Tech Tips





A REAL 64 Core CPU - For SCIENCE!

793,648 views


27K 975 SHARE



Linus Tech Tips  
Published on Oct 31, 2017

SUBSCRIBE 4.6M





A REAL 64 Core CPU - For SCIENCE!

793,648 views


27K 975 SHARE ...



**Linus Tech Tips**  
Published on Oct 31, 2017

SUBSCRIBE 4.6M

Forget Threadripper and Core i9 - Xeon Phi is the go-to for x86 core count. But how does it actually perform?



A REAL 64 Core CPU - For SCIENCE!

793,648 views


27K 975 SHARE ...



**Linus Tech Tips**  
Published on Oct 31, 2017

SUBSCRIBE 4.6M





A REAL 64 Core CPU - For SCIENCE!

793,648 views

27K 975 SHARE ...




Linus Tech Tips

Published on Oct 31, 2017

SUBSCRIBE 4.6M

Forget Threadripper and Core i9. Xeon Phi is the go-to for 48+ cores count. But how does it actually



A REAL 64 Core CPU - For SCIENCE!


793,648 views

27K 975 SHARE ...



**Linus Tech Tips**  
Published on Oct 31, 2017

SUBSCRIBE 4.6M



A REAL 64 Core CPU - For SCIENCE!

793,648 views

27K 975 SHARE



Linus Tech Tips  
Published on Oct 31, 2017

SUBSCRIBE 4.6M

Data  
( $X, y$ )



# Data (X,y)

So that we can train models to predict y  
from X, i.e.


$$X \rightarrow y$$

kaggle

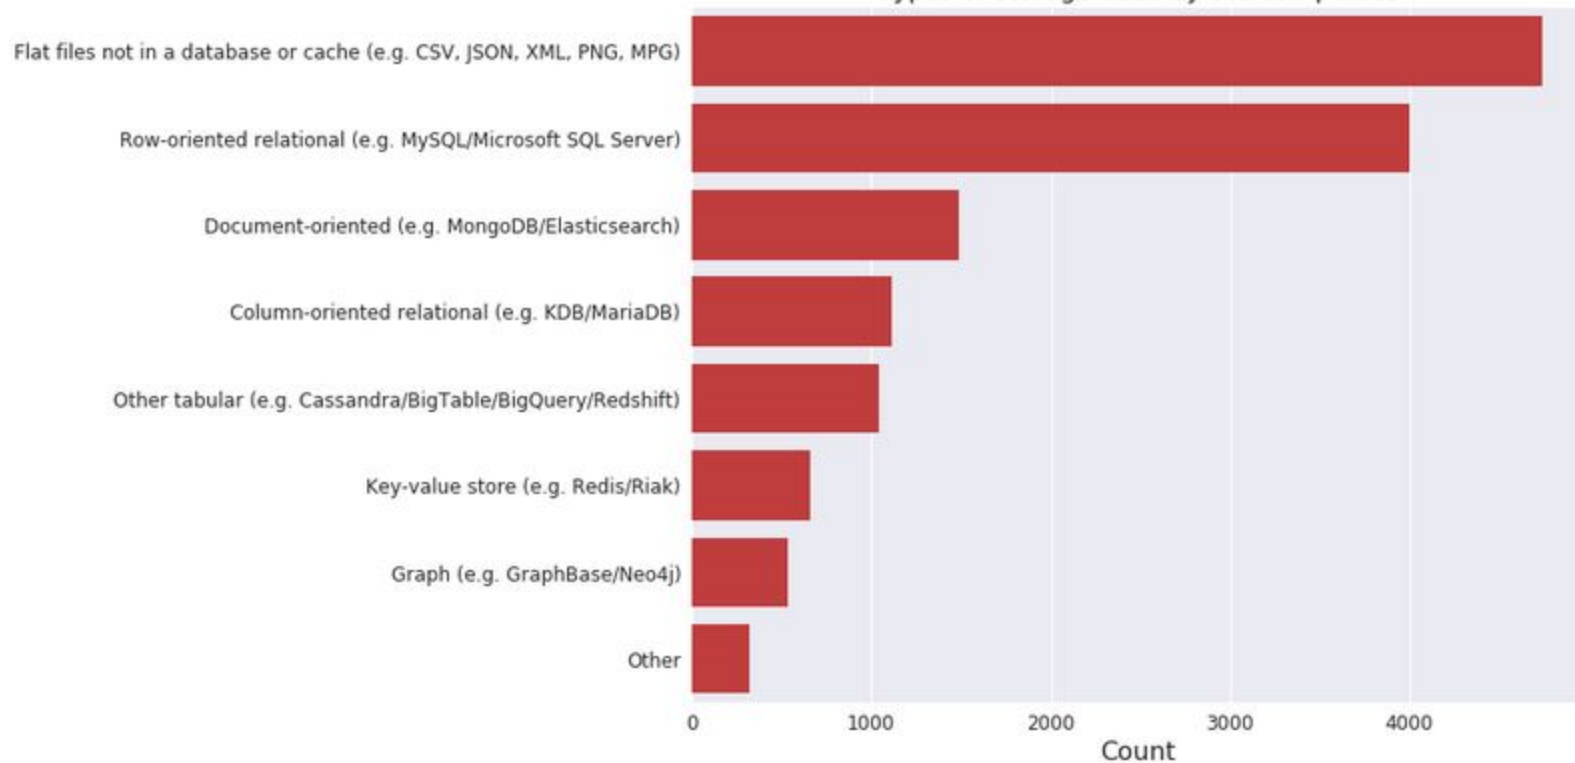
16000 responses

# 2017 The State of Data Science & Machine Learning

“What type of data is used at work?”

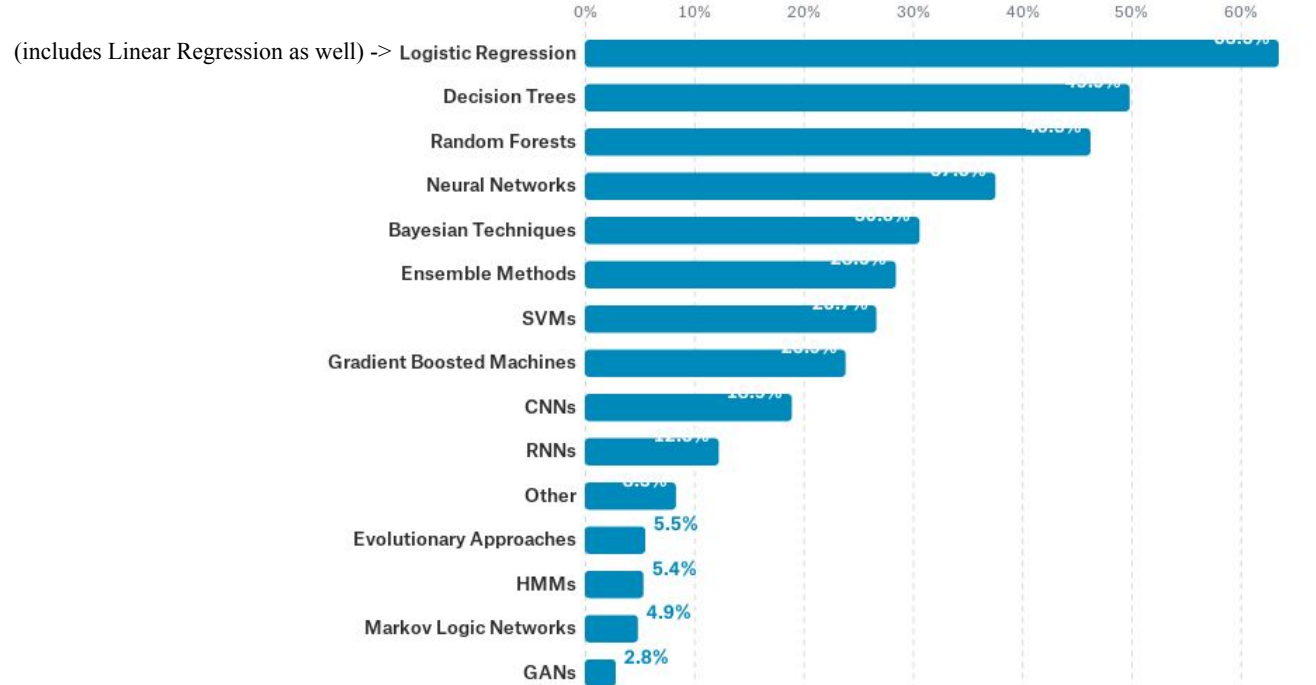


Types of storage used by the companies





## *What people do at work* - Algorithm used at work, selected from multiple choice by responders



7,301 responses


*Machine Learning (ML) methods people want to learn -*

“Which ML, Data Science (DS) method are you most excited about learning in the next year?”

```
print(pd.Series({'count': count, 'percentage': percentage}))
```

Deep learning	: 4362 --> approx. 40.27%
Neural Nets	: 1386 --> approx. 12.79%
Time Series Analysis	: 680 --> approx. 6.28%
Bayesian Methods	: 511 --> approx. 4.72%
Text Mining	: 493 --> approx. 4.55%
Genetic & Evolutionary Algorithms	: 425 --> approx. 3.92%
Social Network Analysis	: 364 --> approx. 3.36%
Anomaly Detection	: 307 --> approx. 2.83%
Ensemble Methods (e.g. boosting, bagging)	: 269 --> approx. 2.48%
Other	: 258 --> approx. 2.38%

“What tools are used at work?”



7,955 responses

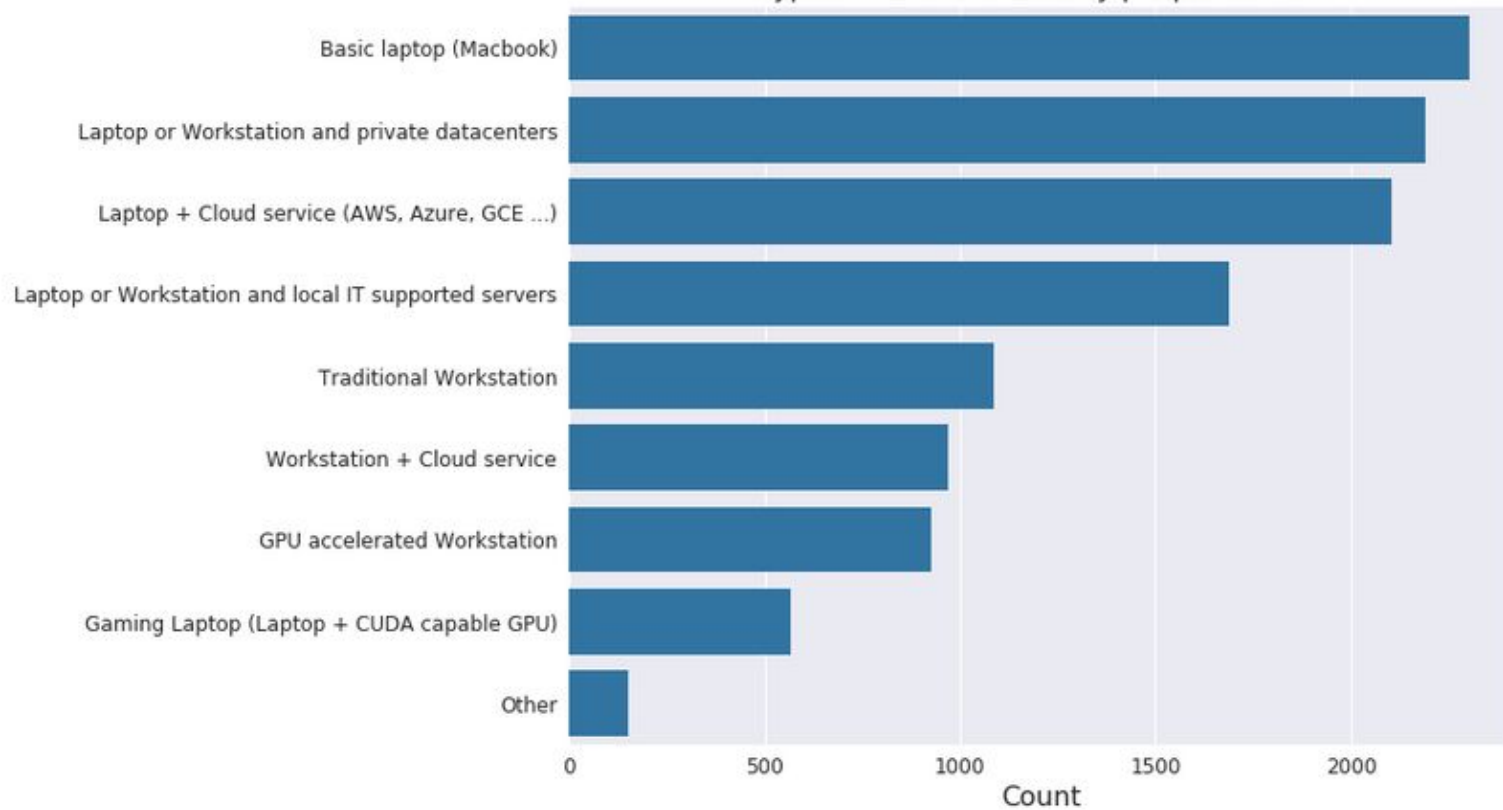
Tools for Next year -  
“Which tool or technology  
are you most excited about  
learning in the next year?”

```
print([tools], [{"tool": choice, "count": count, "percent": percent}])
```

```
TensorFlow : 2621 --> approx. 23.83%  
Python : 1713 --> approx. 15.58%  
R : 910 --> approx. 8.27%  
Spark / MLlib : 755 --> approx. 6.86%  
Hadoop/Hive/Pig : 417 --> approx. 3.79%
```



Type of hardware used by people at work



- GPU(s)
- Deep Learning/Neural Networks (with Linear/Logistic Regression)
- TensorFlow (tf)

# So what is TensorFlow doing now for data pipelines?

Given that Data elements have same type,

*Problem* : Dataset might be too large to materialize all at once ... or infinite

# Functional programming to the rescue!

(Python) Functions on functions, composing functions, functionals

Compose functions like `.map()` and `.filter()` to preprocess

Apply functions such as `.shuffle()` and `.batch()` to transform data to feed into model

# Input pipelines = lazy lists (i.e. lazy evaluation)

```
Python 2.7.13 (default, May 10 2017, 20:04:28)
[GCC 6.3.1 20161221 (Red Hat 6.3.1-1)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> numbers = range(10)
>>> iterator = iter(numbers)
>>> print numbers
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
>>> print iterator
<listiterator object at 0x7f703e5f45d0>
>>> print iterator.next()
0
>>> print iterator.next()
1
>>> print iterator.next()
2
>>> print iterator.next()
3
>>> print iterator.next()
4
>>> █
```



# Introducing `tf.data`

Functional input pipelines in TensorFlow



# The Datasets API

`tf.data`

Better input pipelines for TensorFlow

Or

*“It will revolutionize the way we do data pipelines,  
the way we input data.”*

# Getting started; Install tensorflow-gpu, use virtual-env

```
sudo apt-get install virtual-env (Ubuntu)
```

```
sudo dnf install virtual-env (RedHat/Fedora)
```

```
source ~/tf/bin/activate (go to your virtual environment)
```

```
pip install --upgrade tensorflow-gpu # Python 2.7 and GPU(s) or
```

```
pip3 install --upgrade tensorflow-gpu # Python 3.x and GPU(s)
```

I do not recommend `pip install tensorflow`, for CPU only, as well.

# The Dataset interface

Data sources and functional transformations

`.csv, .txt -> tf.data.Dataset`

# The Dataset interface

Data sources and functional transformations

`.csv, .txt -> tf.data.Dataset`

(X,y)

6.1101	17.592
5.5277	9.1302
8.5186	13.662
7.0032	11.854
5.8598	6.8233
8.3829	11.886
7.4764	4.3483
8.5781	12
6.4862	6.5987
5.0546	3.8166
5.7107	3.2522

(X,y)

120	4	setosa	versicolor	virginica		
6.4	2.8	5.6	2.2	2		
5.0	2.3	3.3	1.0	1		
4.9	2.5	4.5	1.7	2		
4.9	3.1	1.5	0.1	0		
5.7	3.8	1.7	0.3	0		
4.4	3.2	1.3	0.2	0		
5.4	3.4	1.5	0.4	0		
6.9	3.1	5.1	2.3	2		
6.7	3.1	4.4	1.4	1		
5.1	3.7	1.5	0.4	0		
5.2	2.7	3.9	1.4	1		
6.9	3.1	4.9	1.5	1		
5.8	4.0	1.2	0.2	0		
5.4	3.9	1.7	0.4	0		
7.7	3.8	6.7	2.2	2		
6.3	3.3	4.7	1.6	1		
6.8	3.2	5.9	2.3	2		
7.6	3.0	6.6	2.1	2		
6.4	3.2	5.3	2.3	2		

# The Dataset interface

Data sources and functional transformations

`.csv, .txt -> tf.data.Dataset`

`tf.data.TextLineDataset(  
    filename_path  
    (X,y)`

```
6.1101,17.592
5.5277,9.1302
8.5186,13.662
7.0032,11.854
5.8598,6.8233
8.3829,11.886
7.4764,4.3483
8.5781,12
6.4862,6.5987
5.0546,3.8166
5.7107,3.2522
```

(X,y)

```
120,4,setosa,versicolor,virginica
6.4,2.8,5.6,2.2,2
5.0,2.3,3.3,1.0,1
4.9,2.5,4.5,1.7,2
4.9,3.1,1.5,0.1,0
5.7,3.8,1.7,0.3,0
4.4,3.2,1.3,0.2,0
5.4,3.4,1.5,0.4,0
6.9,3.1,5.1,2.3,2
6.7,3.1,4.4,1.4,1
5.1,3.7,1.5,0.4,0
5.2,2.7,3.9,1.4,1
6.9,3.1,4.9,1.5,1
5.8,4.0,1.2,0.2,0
5.4,3.9,1.7,0.4,0
7.7,3.8,6.7,2.2,2
6.3,3.3,4.7,1.6,1
6.8,3.2,5.9,2.3,2
7.6,3.0,6.6,2.1,2
6.4,3.2,5.3,2.3,2
```



```
tf.data.TextLineDataset(file_path)

    .skip(1) # Skip header row

... (.map(...))...

    .batch(m_i) # e.g. m_i = 64

    .repeat(1000) # repeat dataset this number of times, when training
```

```
tf.data.TextLineDataset(file_path)
```

```
.map(decode_csv, num_parallel_calls=m_i) # Skip header row
```

We now have to write custom Python functions to parse data.

```
def decode_csv(line):
```

```
    Parsed_line = tf.decode_csv(line,  
record_defaults=[[0.,],[0.,]], field_delim=",")
```

```
...
```

```
Xandy = parsed_line[:-1], parsed_line[-1]
```

```
return Xandy
```

# The Iterator interface

Sequential access to (batches of) Dataset elements

Create an Iterator from a Dataset:

```
dataset.make_one_shot_iterator()
```

```
dataset.make_initializable_iterator()
```

# The Iterator interface

Sequential access to (batches of) Dataset elements

```
next_X_i, next_y_i = iterator.get_next()
```

```
while ... :
```

```
    X_i, y_i = sess.run([next_X_i, next_y_i])
```

# No more `tf.placeholder`, no more `Queueing`

Use `next_X_i, next_y_i = iterator.get_next()` *directly*.

e.g.

```
yhat = tf.matmul(next_X_i, w) + b
```

# No more `tf.placeholder`, no more Queuing

1 Answer

active

oldest

votes



Placeholder is not needed when using DataSet apis any more, for reading data already a part of

`tf.Graph`.

2



We do not need to read file in python code, and feed them when training, but read data as a tensorflow op in the `tf.Graph`, it will be much more efficiency for tensorflow ops mainly run in cpp.



as in your case, this to lines:

```
x = tf.placeholder(tf.float32, [None, INPUT_SIZE], name='INPUT')
y_ = tf.placeholder(tf.float32, [None, OUTPUT_SIZE], name='OUTPUT')
```

change into:

```
x = next_example
y_ = next_label
```

And **remove** the `feed_dict` when calling `Session.run`

[share](#) [improve this answer](#)

[edited Sep 11 at 13:40](#)

[answered Sep 11 at 13:36](#)



[Tianjin Gu](#)

541 • 2 • 13

cf.

<https://stackoverflow.com/questions/46156778/how-to-map-iterators-output-to-placeholder-in-loss-function-in-tensorflow-using>



```
def input_fn():  
    dataset = ...  
  
    # A one-shot iterator automatically initializes itself on first use.  
    iterator = dataset.make_one_shot_iterator()  
  
    # The return value of get_next() matches the dataset element type.  
    images, labels = iterator.get_next()  
  
    return images, labels  
  
# The input_fn can be used as a regular Estimator input function.  
estimator = tf.estimator.Estimator(...)  
estimator.train(train_input_fn=input_fn, ...)
```

Demo

# tf.data Summary

## **tf.data.Dataset**

Represents input pipeline using functional transformations

## **tf.data.Iterator**

Provides sequential access to (batches of) elements of a **Dataset**

# tf.data Conclusion

Getting your data into TensorFlow with `tf.data`

- Fast
  - `tf.data` is implemented in C++ to avoid Python overhead
- Flexible

*“It will revolutionize the way we do data pipelines, the way we input data.”*





```
Python 2.7.13 (default, May 10 2017, 20:04:28)
[GCC 6.3.1 20161221 (Red Hat 6.3.1-1)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import tensorflow
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/__init__.py", line 24, in <module>
    from tensorflow.python import *
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/python/__init__.py", line 49, in <
module>
    from tensorflow.python import pywrap_tensorflow
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/python/pywrap_tensorflow.py", line
72, in <module>
    raise ImportError(msg)
ImportError: Traceback (most recent call last):
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/python/pywrap_tensorflow.py", line
58, in <module>
    from tensorflow.python.pywrap_tensorflow_internal import *
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/python/pywrap_tensorflow_internal.
py", line 28, in <module>
    _pywrap_tensorflow_internal = swig_import_helper()
  File "/home/topolo/tf/lib/python2.7/site-packages/tensorflow/python/pywrap_tensorflow_internal.
py", line 24, in swig_import_helper
    _mod = imp.load_module('_pywrap_tensorflow_internal', fp, pathname, description)
ImportError: libcublas.so.8.0: cannot open shared object file: No such file or directory
```

Failed to load the native TensorFlow runtime.

See [https://www.tensorflow.org/install/install\\_sources#common\\_installation\\_problems](https://www.tensorflow.org/install/install_sources#common_installation_problems)

for some common reasons and solutions. Include the entire stack trace  
above this error message when asking for help.

```
>>> █
```



# Upgrade to CuDNN 7 and CUDA 9 #12052



tpankaj opened this issue on Aug 4 · 78 comments



tpankaj commented on Aug 4 · edited



## System information

- Have I written custom code (as opposed to using a stock example script provided in TensorFlow): No
- OS Platform and Distribution (e.g., Linux Ubuntu 16.04): Windows Server 2012
- TensorFlow installed from (source or binary): binary
- TensorFlow version (use command below): 1.3.0-rc1
- Python version: 3.5.2
- Bazel version (if compiling from source): N/A
- CUDA/cuDNN version: CUDA V8.0.44, CuDNN 6.0
- GPU model and memory: Nvidia GeForce GTX 1080 TI, 11 GB
- Exact command to reproduce: N/A

## Describe the problem

Please upgrade TensorFlow to support CUDA 9 and CuDNN 7. Nvidia claims this will provide a 2x performance boost on Pascal GPUs.



37



shivaniag added **stat:awaiting tensorflow** **type:feature** labels on Aug 4



shivaniag commented on Aug 4



@tfboyd do you have any comments on this?



tfboyd commented on Aug 4 · edited

Contributor



cuDNN 7 is still in preview mode and is being worked on. We just moved to cuDNN 6.0 with 1.3, which should go final in a couple weeks. You can download cuDNN 1.3.0rc2 if you are interested in that. I have not compiled with cuDNN 7 or CUDA 9 yet. I have heard CUDA 9 is not easy to install on all platforms and

Assign

tfboyd

Labels

type:feature

Project

None yet

Milestone

No milestone

Notification

You're not subscribed to this issue

38 participants

tfboyd

shivaniag

tfboyd

and others



smitshilu commented 2 days ago



I am trying to build 1.4 with CUDA 9 and cuDNN 7 in mac 10.13 high sierra. I am keep getting this error

```
ERROR: /Users/smitshilu/tensorflow/tensorflow/core/kernels/BUILD:2948:1: output 'tensorflow/c
ERROR: /Users/smitshilu/tensorflow/tensorflow/core/kernels/BUILD:2948:1: not all outputs were
Target //tensorflow/tools/pip_package:build_pip_package failed to build
```

Any solution for this?



vellamike commented 2 days ago



@smitshilu possibly related [#2143](#)



grisaitis referenced this issue 2 days ago

**conv2d\_transpose crashing, "NotFoundError: No algorithm worked!", only with batch size  $>= 2^{16}$  #13869**

[Open](#)



ViktorM commented a day ago



Why 1.4 still doesn't have CUDA 9 in binaries? This version was released quit a long ago and for using with V100 building from source is required which is not so smooth and fast following to a number of issues reported.



vellamike commented a day ago



@ViktorM What issues did you have compiling from source? It was a bit tricky but not that hard.



tfboyd commented 20 hours ago

Contributor



26-SEP-2017 was the GA for CUDA 9. If we release CUDA 9 + cuDNN 7 binaries in Q4 I think this will be the fastest we have upgraded cuDNN. I was not here for 8.5 to 9 so I have no idea. I would like us to go a little faster but this also means anyone with a CUDA 8 setup has to upgrade not just to CUDA 9 but they also need to upgrade their device driver to 384.x, which I can say is not something production people take lightly.

Ideally we would have infinite (or just a few more but the matrix explodes fast) builds, but that is another problem that would take a long time to explain and I doubt many people care.



**yaroslavvb** commented 14 hours ago

Contributor



BTW, I [observed](#) 85 T ops/second with CUDA 9 on float16 matmul/V100 using Nvidia's [NGC](#) TensorFlow container (as opposed to 8.8 T ops/s on my GTX 1080 at home). Really looking forward to having these improvements in the officially supported version!



**tfboyd** commented 14 hours ago

Contributor



**@yaroslavvb** Being very honest we are working through some FP16 issues. There is a path in the `tf_cnn_benchmarks` for FP16 and the focus is on ResNet50 first and we are working on auto scaling for FP16 as well. You can give it a try if you are interested but we are actively working through some problems. People are on it and it is just taking time. We finally have DGX-1s in house so we can also play with the same containers and try to keep track of performance on that exact platform moving forward.



**vickylance** commented 17 minutes ago



Ok, So I am going to install Ubuntu 17.10 and I just wanted to try all the latest stuffs for fun. Before I do I just wanted to know did any one try the below stacks building from source and got any luck?

-> Ubuntu 17.10, CUDA 9.0, cuDNN 7.0, TF master

-> Ubuntu 17.10, CUDA 8.0, cuDNN 6.1, TF 1.4



**aluo-x** commented 7 minutes ago



I am encountering the same issue as **@xsr-ai**, specifically using Python 3.6.3, VS 2017, CUDA 9, cuDNN 7.



# cuBlackDreams

<https://github.com/ernestyalumni/cuBlackDream>

- CUDA C++14, CUDA 9, C++14 smart pointers
- Parallel programming, parallel algorithms only
- Fast File I/O with C++14 binary format (`std::ios::binary`)
- Multi-GPU ready (using CUDA Unified Memory Management)

Machine Learning and Deep Learning in CUDA, for the GPU

```
topolo@localhost:~/PropD/cuBlackDream/examples
File Edit View Search Terminal Help
activationf.o  linreg.exe  logreg.exe  RModule.exe  testAxon_act.cu
Axon.o         LinReg.ipynb LogReg.ipynb RModule.ipynb testAxon_act.exe
FileIO.o       linreg.o    Makefile    RModule.o
linreg.cu      logreg.cu   RModule.cu  smartCUBLAS.o
[topolo@localhost examples]$ ./linreg.exe
d : 1
K : 1
m : 97

X_exldataout.size() : 97
SIZE X = h Xvec.size() : 97
result of linReg cost : 32.0727
Time to grad_desc 1500 iterations : 2396.3 ms
hTheta1 : 1.16641
hb1 : -3.63077

multi-dim. linear reg case :
d : 2
K : 1
m : 47

X_exldata2out.size() : 94
Time to multi-dim. grad_desc 400 iterations : 668.1 ms
hTheta (multi) : 117935 68621.7
hb (multi) : 315585
[topolo@localhost examples]$ pwd
/home/topolo/PropD/cuBlackDream/examples
[topolo@localhost examples]$
```

<https://github.com/ernestyalumni/cuBlackDream>

```
/home/topolo/PropD/cuBlackDream/examples
[topolo@localhost examples]$ ./logreg.exe

For ex2data1.txt :
d = 2, K = 1, m = 100
costJ for cross-entropy function, at initial theta (zeros) : 0.693147
Expected cost (approx): 0.693

Cost at test theta: 0.21833
Expected cost (approx): 0.218
Time to grad_desc 1500 iterations : 2845714.2 ms
Cost at theta found by grad desc: 0.21542
Expected cost (approx): 0.203
hTheta1 : 0.144729 0.139161
hb1 : -17.4615
[topolo@localhost examples]$
```

<https://github.com/ernestyalumni/cuBlackDream>